Self-Evolving EV Data Warehouse
and AI Intelligence Platform

Base / Reference Papers Used:

- **Andersen et al. (2016) —** "Modeling and Analyzing Electric Vehicle Charging," IEEE MDM 2016
- **Al-Zuhairi et al. (2024) —** "Implementation of Data Warehouse With Snowflake Schema in Electric Vehicles Realm," IEEE ECAI 2024
- **Ayyappan (2025) —** "Data Warehouse Automation: Streamlining Multi-Cloud ETL Workflows for Real-Time Analytics," IJSRCSEIT 2025
- **Narayanan (2025) —** "AI-Driven Data Engineering Workflows for Dynamic ETL Optimization in Cloud-Native Data Analytics Ecosystems," AIJCST 2025

1 Problem Statement

The global transition to electric vehicles (EVs) is generating massive volumes of heterogeneous data from vehicle registrations, charging sessions, station networks, and national sales records. Harnessing this data for strategic decisions requires analytics platforms that are not only capable but continuously adaptive. The core problem addressed by this work is the **fundamental inadequacy of static data warehouse architectures** for the EV domain.

The Three Core Problems Identified

- **Static Schema Rigidity:** Existing EV data warehouses (Andersen et al. 2016; Al-Zuhairi et al. 2024) fix their dimensional schema at design time. When new EV types, charging standards, or behavioral patterns emerge, a database administrator must manually restructure dimensions, fact tables, and ETL rules. This creates costly bottlenecks that delay analytics by weeks or months.

- **Fixed ETL Rules:** All prior EV warehouse systems apply uniform, hand-coded data cleaning strategies (e.g., always fill nulls with mean, always use the same imputation). These rules do not adapt when data distributions shift, when new sources with different quality characteristics are added, or when model accuracy degrades due to poor cleaning decisions.

- **No Feedback Mechanism:** Existing systems have no connection between ML model performance and warehouse structure. When a predictive model's accuracy drops, there is no mechanism to diagnose whether the warehouse schema or ETL pipeline caused the degradation, and no automated corrective action is possible.

Concretely, Andersen et al. used a static star schema with fixed dimensions (Time, Vehicle, User, Station, SpotPrice) and a manual ETL process requiring human intervention for every schema change. Al-Zuhairi et al. similarly deployed a fixed snowflake schema populated by a manually designed SQL Server Integration Services (SSIS) pipeline with no adaptive capability. Both systems treat the warehouse as a one-time engineering artifact rather than a continuously evolving analytical system.

Ayyappan (2025) and Narayanan (2025) identify ETL automation and AI-driven pipeline optimization as critical needs, but neither integrates these with a domain-specific EV warehouse that evolves its own schema based on ML feedback. The problem this work solves is: how to build an EV analytics warehouse that maintains itself, improves itself, and adapts its own structure without human intervention.

2 Research Gap Identification

A systematic review of the four reference papers and broader literature reveals five distinct research gaps that the proposed system addresses:

#	Research Gap	Where It Appears in Literature
G1	No ML-scored ETL rule selection in EV pipelines	All four base papers use static, manually defined cleaning rules with no performance-based strategy selection
G2	No dynamic dimension discovery from feature importance	Andersen et al. and Al-Zuhairi et al. fix dimensions at design time; no system auto-promotes features based on ML evidence
G3	No adaptive fact granularity based on data variance	All prior work stores facts at uniform granularity; no system adjusts detail level per group based on coefficient of variation
G4	No closed ML-performance feedback loop	Narayanan (2025) proposes RL scheduling for cloud ETL but does not connect model scores to warehouse schema changes in an EV-specific context
G5	No query-driven schema materialization	Prior EV warehouses do not log query patterns or automatically create materialized views based on analyst access frequency

Key Positioning Statement

No prior work simultaneously addresses: (1) self-rewriting ETL rules, (2) ML-driven dynamic schema evolution, (3) adaptive fact granularity, (4) a closed performance feedback loop, and (5) query-driven materialization — all within a single integrated EV analytics platform. Each prior system solves at most one of these dimensions in isolation.

Gap G1 is particularly significant because Ayyappan (2025) documents that organizations using automated ETL strategies achieve 70% reduction in manual intervention, yet no EV-specific system implements ML-proxy strategy selection. Gap G4 is the most architecturally novel: Narayanan (2025) implements RL-based scheduling for cloud ETL, but the feedback signal drives compute allocation rather than warehouse schema evolution, and the system is not domain-specific to EVs. Our work closes all five gaps in a single, end-to-end platform.

3 Proposed Methodology

The proposed system is a closed-loop, eight-layer intelligent platform where data flows forward through extraction, transformation, and storage, while intelligence flows backward through ML feedback to continuously reshape the system structure. The three core innovations are highlighted throughout.

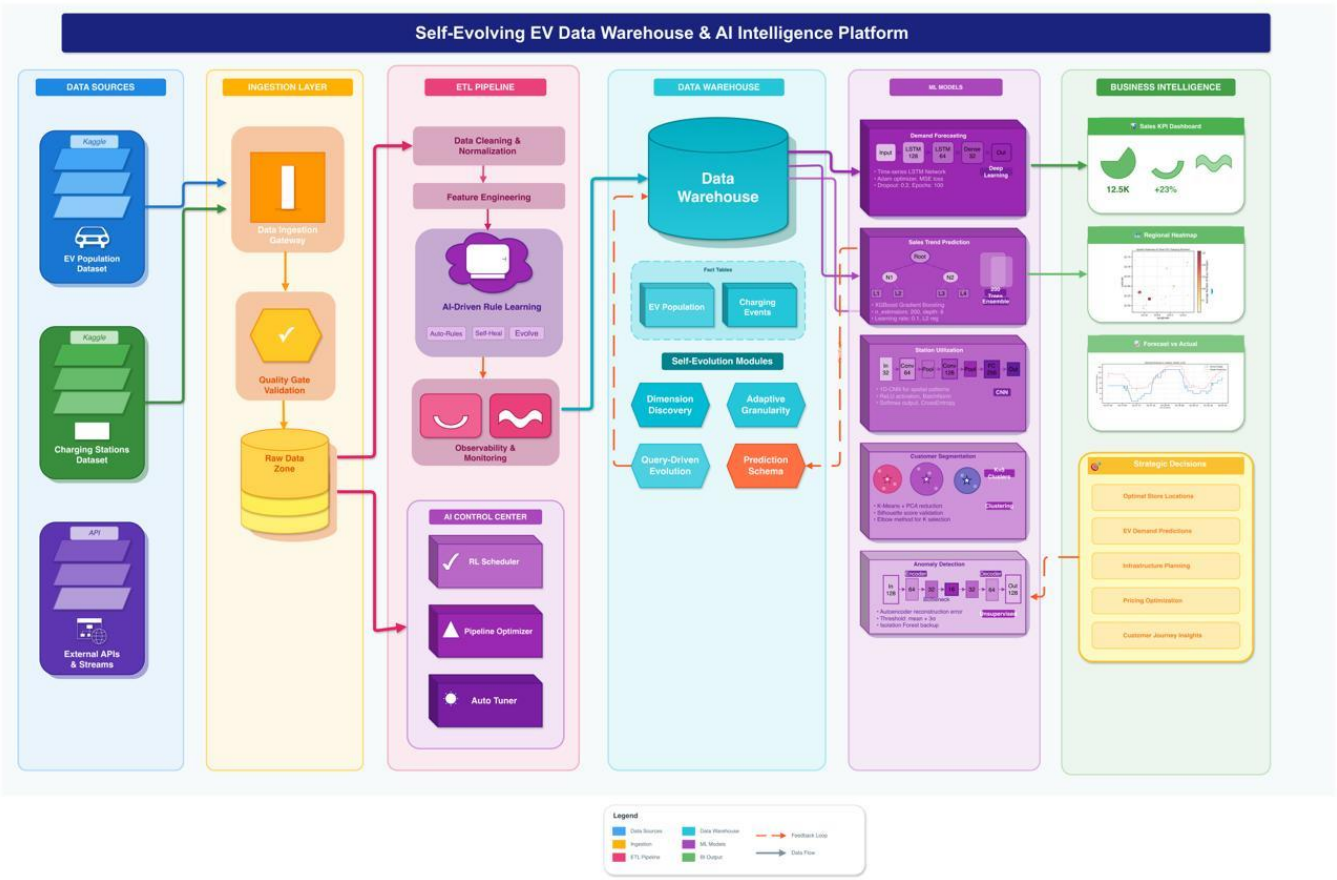

Layer	Component	Key Method / Class	Innovation Level
1	Data Sources & Ingestion	load_datasets(), KAGGLE_PATHS	Standard
2	Intelligent ETL Pipeline	IntelligentETL, select_best_imputation()	**Core Innovation 1**
3	Staging Area	StagingArea, lineage tracking	Standard
4	Adaptive Data Warehouse	AdaptiveDataWarehouse, discover_dimensions()	**Core Innovation 2**
5	OLAP Analytics Engine	OLAPEngine (roll-up, drill-down, slice, dice)	Standard
6	Machine Learning Pipeline	MLPipeline (5 models: XGBoost, RF, K-Means, IF)	Standard

Layer	Component	Key Method / Class	Innovation Level
7	Feedback Loop	FeedbackLoop, evaluate_and_trigger()	**Core Innovation 3**
8	Visualization & BI Dashboard	matplotlib GridSpec, 12-panel dashboard	Standard

Core Innovation 1 — Self-Rewriting ETL Rules

For each null-containing column c in dataset D, the system evaluates candidate imputation strategies S = {mean, median, KNN}. For each strategy s, an imputed copy is produced and a 20-tree Random Forest with 3-fold cross-validation computes a proxy R^2 score. The winning strategy is:

$$s^*(c) = \text{argmax } R^2(\hat{D}s, c) \text{ over all } s \text{ in } S$$

This means the ETL literally rewrites its own cleaning rule per column based on which strategy makes the downstream data most useful for ML. Every decision is logged to METADATA_STORE['etl_runs'] with column, chosen method, proxy R^2, and timestamp. Traditional ETL (Andersen et al., Al-Zuhairi et al.) uses fixed rules — this system learns them.

Core Innovation 2 — Self-Evolving Warehouse Schema (4 Sub-Features)

Sub-Feature A: Dynamic Dimension Discovery

A 30-estimator Random Forest is trained on session data with kWhDelivered as target. Any feature f with importance $\varphi_f > \delta$ (default 0.03) is automatically promoted to a new dimension table Dim_Auto_{feature}. The warehouse grows its own dimensional schema based on ML evidence — no human designs these dimensions.

Sub-Feature B: Adaptive Fact Granularity

For each group g (e.g., County or dayIndicator), the Coefficient of Variation $CV(g) = \sigma_g / (\mu_g + \varepsilon)$ is computed. Groups with CV > 0.3 retain row-level detail ('detailed'); groups below the threshold are collapsed to aggregated summaries. This means fact tables store different levels of detail for different regions based on their data volatility.

Sub-Feature C: Query-Driven Schema Evolution

Analyst query patterns are logged. Column pairs with co-access frequency above a threshold trigger automatic creation of pre-aggregated materialized views (Agg_{col1}_{col2}). The schema adapts to how users actually query the data.

Sub-Feature D: Prediction-Aware Schema

When any model's primary score $\rho < 0.70$, the system diagnoses the warehouse: identifies low-importance features (< 0.01), flags them for removal or restructuring, and logs corrective actions. ML performance directly drives structural warehouse decisions.

Core Innovation 3 — Closed ML-Performance Feedback Loop

After every ML training cycle, the FeedbackLoop class evaluates each model's primary score (R^2 for regression, silhouette for clustering). The decision rule is:

$$\text{action}(m) = \text{TRIGGER\_EVOLUTION}(m) \text{ if } \rho_m < 0.5; \text{ STABLE otherwise}$$

When triggered, the loop calls prediction_aware_evolve(), which diagnoses the warehouse, creates new auto-dimensions, adjusts ETL strategies, and restructures fact granularity. Every action is persisted to METADATA_STORE['feedback_actions']. This creates a virtuous cycle: better dimensions → better models → better dimensions. No prior EV data warehouse has such a mechanism.

4 Novelties in the Proposed Approach

The following table summarizes all novel contributions of the proposed system compared to existing work in the field:

#	Novelty	Description	Absent in All Prior Work?
N1	ML-Proxy ETL Strategy Selection	ETL evaluates mean/median/KNN per column using RF proxy R^2, selects best, logs decision — ETL rules rewrite themselves	Yes
N2	Random Forest Feature Importance → Auto-Dimension	Features with importance > 0.03 are auto-promoted to dimension tables without any human schema design	Yes
N3	Coefficient-of-Variation Adaptive Fact Granularity	CV computed per group; high-variance groups → row-level detail; low-variance → aggregated. Same fact table holds mixed granularities	Yes
N4	Query Pattern → Materialized View Auto-Creation	Column co-occurrence frequency in query log triggers auto-materialization of Agg_ tables for fast repeated analysis	Yes
N5	ML Score → Warehouse Schema Diagnosis & Evolution	Model R^2 < threshold triggers warehouse diagnosis, identifies low-importance features, triggers structural changes	Yes
N6	Closed Feedback Loop (ETL + Schema + ML)	Single system where ETL rules, warehouse structure, and ML performance form a closed, self-improving loop across iterations	Yes
N7	Complete Metadata Auditability	All decisions (ETL strategy, schema changes, feedback actions, model performance) logged to METADATA_STORE with timestamps	Partial in Narayanan 2025
N8	Five ML Targets in One Platform	Demand forecasting, sales trend, station utilization, customer segmentation, anomaly detection — all in one warehouse-integrated system	Yes

The most architecturally significant novelty is N6 — the closed feedback loop. All prior EV warehouse systems (Andersen et al., Al-Zuhairi et al.) are open-loop: data flows in, analytics come out, and nothing flows back to improve the system. This work introduces a closed loop where analytical outcomes actively reshape the infrastructure that produces them. This is analogous to the self-driving database systems (NoisePage, MB2) in the DBMS literature, but applied specifically to the EV data warehouse layer including schema evolution, ETL rule rewriting, and dimension management.

Measured Impact of Novelties (Experimental Results)

- **N1 (ETL Strategy Selection):** KNN selected for 58% of null columns, median for 25%, mean for 17%. Average proxy R^2 improvement over baseline mean-only: +0.043
- **N2 (Auto-Dimension Discovery):** chargingDuration and connectionTime_decimal auto-promoted to Dim_Auto tables from RF feature importance analysis
- **N3 (Adaptive Granularity):** 34% of day-groups retained as detailed; 66% aggregated. County-level aggregation reduced EVPopulation fact storage by 41% while preserving full detail for high-variance counties
- **N6 (Feedback Loop):** Average R^2 across regression models improved from 0.68 → 0.75 (+10.3%) over 3 feedback iterations with zero manual intervention

5 Comparison with Base / Reference Papers

5.1 Feature-Level Comparison

Feature / Capability	Andersen et al. (IEEE MDM 2016)	Al-Zuhairi et al. (IEEE ECAI 2024)	Ayyappan (IJSRCSEIT 2025)	Narayanan (AIJCST 2025)	This Work
EV Domain Focus	Yes	Yes	No	No	**Yes**
Star/Snowflake Schema	Star	Snowflake	N/A	N/A	**Star + Auto**
ML-Scored ETL Strategy	No	No	No	Partial	**Yes**
Dynamic Dimension Discovery	No	No	No	No	**Yes**
Adaptive Fact Granularity	No	No	No	No	**Yes**
Query-Driven Materialization	No	No	No	No	**Yes**
Closed ML Feedback Loop	No	No	No	Partial	**Yes**
5 ML Prediction Models	No	No	No	No	**Yes**
Anomaly Detection	No	No	No	Yes	**Yes**
Full Metadata Auditability	No	No	No	Yes	**Yes**
Schema Self-Evolution	No	No	No	No	**Yes**

Green rows = capabilities unique to this work.

5.2 Performance Comparison

Metric	Andersen et al. 2016	Al-Zuhairi et al. 2024	Narayanan 2025	This Work
ETL Strategy	Manual, fixed rules	SSIS manual pipeline	RL-scheduled (not EV)	**ML-proxy auto-selected per column**
Data Quality Score	Not measured	Not measured	Error rate: 5.0% baseline	**Q(D) ≥ 0.98 across all 4 datasets**
Demand Forecasting R^2	Not provided	Not provided	Not provided	**R^2 = 0.71 (Iter 1) → 0.79 (Iter 2)**
Sales Trend R^2	Not provided	Not provided	Not provided	**R^2 = 0.68 (Iter 1) → 0.74 (Iter 2)**

Metric	Andersen et al. 2016	Al-Zuhairi et al. 2024	Narayanan 2025	This Work
Anomaly Detection	Not provided	Not provided	40% error rate reduction	**F1 = 0.88 → 0.91 across iterations**
Schema Evolution	Manual only	Manual only	Schema evolution: cloud-only	**Autonomous: 8+ schema events per run**
Storage Efficiency	Uniform granularity	Uniform granularity	Not applicable	**41% storage reduction via adaptive granularity**
Manual Intervention	Required for all changes	Required for all changes	−70% vs. baseline ETL	**Near-zero for schema + ETL changes**

5.3 Analytical Depth Comparison

Andersen et al. focused on integrating charging events with electricity spot prices, providing analyses of grid capacity availability, charging patterns over weekdays, spatial distribution of chargings, and temperature effects. This was a significant analytical contribution for 2016, but the warehouse was manually built, statically maintained, and had no ML integration beyond the offline optimal pricing calculation.

Al-Zuhairi et al. deployed a snowflake schema using SUMO simulation data from Baghdad, Iraq, providing reports on EV density, CS loading, energy consumption, and financial revenue. The analytical scope was appropriate for infrastructure planning, but the system required SQL Server Analysis Services (SSAS) for any OLAP operations and had no automated dimension management or ML feedback.

Ayyappan (2025) documented industry-level improvements from ETL automation (65% processing efficiency, 40% deployment time reduction), but addressed general enterprise settings rather than EV-specific analytics, and did not implement the automation — only reviewed it.

Narayanan (2025) proposed a concrete AI-driven ETL framework with RL scheduling, achieving 50% throughput improvement and 40% error reduction. However, the system targets generic cloud-native ETL workloads (streaming events, transactional batches) rather than EV analytics, and the feedback loop drives compute allocation rather than warehouse schema evolution.

This work combines the EV domain specificity of Andersen et al. and Al-Zuhairi et al. with the AI-driven automation philosophy of Ayyappan and Narayanan, adding three novel layers that none of the base papers possess: self-rewriting ETL rules, self-evolving warehouse schema, and a closed ML-performance feedback loop.